

Build Phases & AI Learning Roadmap

From weekend CLI → community tool → paid SaaS → AI-for-SMB capability

Owner: Srikanth Pusapati · Version 1.0 · July 31, 2026

Part A — Build Phases

Each phase ends with something you can put in front of real people. Budget assumes 5–15 hrs/week. Dates assume a start the week of Aug 3, 2026.

Phase	When	Build	Deliver to whom	Exit criteria
0. Setup	Aug 3 (1 evening)	GitHub repo ChangeLogScriptGenerator, push docs/ (these files), uv project skeleton	—	Repo live, CLI --help runs
1. CLI MVP	Aug 3–16	Collector → classifier → LLM → Markdown (steps in Design doc §7)	Yourself + 2 dev friends' repos	Publishable changelog from a real repo in <2 min
2. GitHub Action + polish	Aug 17–30	action.yml, README with GIF demo, PyPI release, sample outputs	10 dev friends, your team at work, 1–2 Discords	5 people ran it; 5 honest complaints collected
3. Public launch + validation	Aug 31–Sep 13	Landing page w/ email capture, Show HN + r/webdev posts, GitHub Marketplace listing	Communities	Gate: ≥20 signups or 10 active repos, else archive w/ retro
4. Hosted Pro	Sep–Oct	GitHub App, dashboard, hosted changelog pages, Stripe (\$15/29)	Waitlist from Phase 3	First paying team
5. Compound	Nov+	Slack/email digests, brand voice, Jira/Linear ingest, approval flow	Paying users' requests only	\$500 MRR → graduate out of PersonalProjects

Weekly operating rhythm

- 2 build sessions (2–3 hrs each, calendar-blocked) + 1 distribution hour (show someone, post somewhere, answer users) + 30 min learning (Part B). Distribution hour is not optional — it is the difference between a repo and a product.

Part B — AI/LLM Learning Roadmap

Structured so every topic is learned by shipping a piece of this project, then generalized. Order matters — each stage unlocks the next.

Stage 1 — LLM engineering fundamentals (Aug, alongside Phases 1–2)

- **Prompt engineering:** system vs user prompts, few-shot examples, tone control. Practice: your `release.j2` template. Read: Anthropic prompt-engineering docs (docs.claude.com), OpenAI cookbook.
- **Structured outputs:** JSON schema constraints, validation, retry-with-feedback. Practice: `llm.py`. This single skill separates demo-builders from production-builders.
- **Token economics:** counting, batching, model-tier selection. Practice: your $< \$0.05/\text{release}$ budget.
- **Evals:** build the 10-case eval file from Design doc §8; learn LLM-as-judge grading. Read: Anthropic evals guide; look at `promptfoo` as tooling.

Stage 2 — Agentic patterns (Sep–Oct, alongside Phases 3–4)

- **Tool use / function calling:** let the model call functions (e.g., `'fetch_pr_details'`) instead of pre-feeding everything. Practice: upgrade the classifier to request extra context only when needed.
- **Agent loops:** plan → act → observe → repeat; when agents beat pipelines (rarely!) and how to bound them. Read: Anthropic's 'Building effective agents' essay — the pipeline-first philosophy in your design doc comes from it.
- **MCP (Model Context Protocol):** build a tiny MCP server exposing `'generate_changelog'` as a tool — instantly usable from Claude Desktop/Code. Great demo for dev friends and a marketable skill on its own.
- **Claude Agent SDK / frameworks:** study one framework deeply (Claude Agent SDK or LangGraph) by building the Phase-4 'draft announcement + await approval + publish' flow as a small stateful agent.

Stage 3 — Productizing AI for SMBs (Oct+, alongside Phase 5)

- **RAG basics:** embeddings, chunking, retrieval. Practice: 'ask your changelog' search box on hosted pages — small, real, demoable.
- **Ops for AI products:** logging LLM calls, cost dashboards, fallbacks, versioned prompts as deploys. Practice: hosted Pro's admin page.
- **Packaging AI as a service:** study productized-service models (fixed-scope AI automations for small businesses: intake bots, report generators, back-office pipelines). Your changelog tool is the portfolio proof.

Part C — The Transformation Path (→ AI SMB)

The end-state is not 'one changelog tool'; it is you as a person who can reliably turn a business's toil into an AI-powered product. The route:

- **Ship one real AI product end-to-end** (this project): pipeline design, structured outputs, evals, cost control, distribution, billing. This vocabulary + repo is your credibility asset.
- **Teach in public:** write 3–5 short posts as you build ('How I keep my LLM under \$0.05/release', 'Structured outputs saved my changelog bot'). Post where you launched. This compounds into inbound interest — the community delivery you asked for.
- **Do 1–2 paid AI automations for people you know:** a friend's business, a local company — fixed scope, fixed price (\$1–5K): an intake summarizer, a quote generator, a report writer. Reuse the exact pipeline skeleton from this project.
- **Choose the compounding lane by December:** (a) grow the SaaS if MRR is moving, (b) productized AI-automation studio for SMBs if service demand is stronger, or (c) both with the SaaS as lead magnet. Decide with data from steps 1–3, and record the decision in PersonalProjects.

Curated resource list

- Anthropic: docs.claude.com (prompting, tool use, Agent SDK), 'Building Effective Agents' essay, Anthropic cookbook on GitHub.
- OpenAI Cookbook (structured outputs, evals patterns) — read for patterns even if you ship on Claude.
- MCP: modelcontextprotocol.io quickstart — build your first server in an evening.
- Evals tooling: promptfoo (open source) — wire your eval file into it in Stage 1.
- Communities: Hacker News, r/SideProject + r/webdev, Indie Hackers, MicroConf Connect, local AI meetups (great for finding SMB automation clients).
- Business of micro-SaaS: 'The SaaS Playbook' (Rob Walling), Indie Hackers interviews of changelog/dev-tool founders.